

Geometry of the Affective Brain: Manifold Structure, Representational Similarity, and Subject-Independent Decoding of 27 Fine-Grained Emotions from Scalp EEG

Dhruv Kithany

2026-09-07

```
import sys
import json
from pathlib import Path

# Setup paths relative to memo location
memoDir = Path(".").resolve()
projectRoot = memoDir.parent.parent
if str(projectRoot) not in sys.path:
    sys.path.insert(0, str(projectRoot))

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Set scientific visualization defaults
sns.set_theme(style="whitegrid", palette="deep", font_scale=0.95)

# Load pre-computed benchmark and manifold metrics
figuresDir = projectRoot / "figures"
dataDir = projectRoot / "data"

with open(figuresDir / "02_manifold_metrics.json", "r") as f:
    manifoldMetrics = json.load(f)

with open(figuresDir / "03_benchmark_results.json", "r") as f:
    benchmarkMetrics = json.load(f)

with open(figuresDir / "04_generative_metrics.json", "r") as f:
    generativeMetrics = json.load(f)

# Load participant summary
dfParticipants = pd.read_csv(dataDir / "participants_info.csv")
nParticipants = len(dfParticipants)
nMales = int((dfParticipants["Gender"] == 1).sum())
nFemales = int((dfParticipants["Gender"] == 2).sum())

# Extract key statistics
twoNnDim = manifoldMetrics["intrinsic_dimensionality"]["two_nn_dimension"]
pcaVar80 = manifoldMetrics["intrinsic_dimensionality"]["pca_var_80pct"]
pcaVar95 = manifoldMetrics["intrinsic_dimensionality"]["pca_var_95pct"]
rsaR = manifoldMetrics["rsa"]["pearson_r"]
rsaP = manifoldMetrics["rsa"]["pearson_p"]
rsaRho = manifoldMetrics["rsa"]["spearman_rho"]

# Benchmark top performances
rfQuadAcc = benchmarkMetrics["task_a_quadrants"][0]["accuracy"] * 100
rfQuadF1 = benchmarkMetrics["task_a_quadrants"][0]["macro_f1"] * 100
mlpQuadAcc = benchmarkMetrics["task_a_quadrants"][1]["accuracy"] * 100

rfCowenTop1 = benchmarkMetrics["task_b_cowen27"][0]["accuracy"] * 100
```

```
rfCowenTop5 = benchmarkMetrics["task_b_cowen27"][0]["top_5_accuracy"] * 100
mlpCowenTop1 = benchmarkMetrics["task_b_cowen27"][1]["accuracy"] * 100
mlpCowenTop5 = benchmarkMetrics["task_b_cowen27"][1]["top_5_accuracy"] * 100
```

The Problem: Mapping Brain Waves as High-Dimensional Geometry

In machine learning and computational geometry, one big question is whether continuous biological signals capture complex human experiences like emotion. Most emotion classification papers treat brain signals as a simple 2-class or 4-class problem (like “happy vs. sad”). But human feelings are much richer than that.

Cowen and Keltner (2017) showed through massive self-report studies that human emotion is organized into at least **27 distinct categories** along a continuous semantic geometry.

Here, we take raw scalp EEG data (14 electrodes, 128 Hz) across 88 participants and extract a **462-dimensional feature space** ($D = 462$) per trial. From a computer science and geometric ML standpoint, we want to answer three specific questions:

1. **Topological Alignment & Intrinsic Dimension:** Does high-dimensional brain wave space actually reflect the geometry of psychological emotion models (Russell 1980) when we run Representational Similarity Analysis (RSA) (Kriegeskorte, Mur, and Bandettini 2008)? And if we measure 462 features, what is the *true* non-linear intrinsic dimensionality of the underlying manifold (Facco et al. 2017)?
2. **Subject-Independent Decoding:** Can standard ML classifiers and neural networks decode 27 fine-grained emotions under a strict **Leave-Subjects-Out (LSO)** split where the test participants are completely unseen during training?
3. **Generative Modeling with Covariance Shrinkage:** Because $D = 462$ is large relative to the number of trials per class, sample covariance matrices collapse and become singular. Can class-conditional Gaussian models with Ledoit-Wolf shrinkage (Ledoit and Wolf 2004) synthesize realistic, high-fidelity synthetic brain trials?

Executive Summary of Results

We analyzed the benchmark EEGEmotions-27 dataset (Phuong et al. 2025) (11,235 trials recorded from 88 subjects: 43 males, 45 females, healthy young adult university cohort). Here is what the algorithms and geometry revealed:

1. **The 462D Space Collapses to ~13 Dimensions:** While we measure 462 features per trial, the Two-NN algorithm (Facco et al. 2017) proves the intrinsic dimensionality of the neural emotion manifold is actually $d \approx 13.4$. That means ~97% of the ambient dimensions are redundant or correlated noise. Furthermore, our Representational Similarity Analysis (RSA) between the 27x27 EEG centroid distance matrix and the psychological Valence-Arousal distance matrix shows statistically significant geometric alignment ($r = 0.142$, $p = 0.0077$; Spearman rank $\rho = 0.143$, $p = 0.0073$).
2. **Decoding Without Identity Leakage:** In a strict Leave-Subjects-Out benchmark (74 training subjects, 14 unseen holdout subjects):
 - **4-Class Circumplex Quadrants:** Random Forest hits **56.64% accuracy** (vs. 25.0% random chance) and 43.40% Macro F1, followed by a Bottleneck MLP at 50.20%.
 - **27 Fine-Grained Emotions:** Random Forest achieves **30.83% Top-1 accuracy** and **79.83% Top-5 accuracy** (vs. 3.70% random chance), with our deep Bottleneck MLP achieving 80.28% Top-5 accuracy.
3. **Biological Sanity Check:** Frontal Alpha Asymmetry (FAA: $\ln(AF4) - \ln(AF3)$) cleanly separates approach/positive states (Admiration, Joy, Excitement) from avoidance/negative

states (Fear, Sadness, Disgust), confirming our feature pipeline captures real cognitive signals rather than noise (Davidson 1992).

4. **Synthetic Data Generation:** Ledoit-Wolf regularized covariance matrices allow us to sample synthetic trials in 462D space without matrix singularity, reaching a strong Fréchet Distance of $d_F = 116.88$ for Sadness and $d_F = 401.60$ for Anger.

Part 1: Signal Processing & Frontal Asymmetry

Electrode Layout and Signal Processing

Scalp EEG measures aggregate electrical potentials from the brain using electrodes placed on the head. In this dataset, participants wore an Emotiv EPOC+ headset with 14 channels arranged according to the standard international 10-20 grid: Prefrontal (AF3, AF4), Frontal (F7, F3, FC5, FC6, F4, F8), Temporal (T7, T8), Parietal (P7, P8), and Occipital (O1, O2).

To extract features from each trial’s time-series data: 1. We compute Welch’s Power Spectral Density (PSD) using overlapping Hann windows, calculating total energy across 5 standard frequency bands: Delta (1–4 Hz), Theta (4–8 Hz), Alpha (8–13 Hz), Beta (13–30 Hz), and Gamma (30–45 Hz). 2. We compute nonlinear time-domain statistics including Hjorth parameters (Activity, Mobility, Complexity) and Shannon entropy. 3. We compute hemispheric log-power differences across symmetric channel pairs, yielding a total feature vector of $D = 462$ dimensions per trial.

```
fig, axes = plt.subplots(2, 1, figsize=(11, 4.8), dpi=200)

imgJoy = plt.imread(figuresDir / "01_topomap_multiband_joy_hvha.png")
imgSad = plt.imread(figuresDir / "01_topomap_multiband_sadness_lvla.png")

axes[0].imshow(imgJoy)
axes[0].axis("off")
axes[0].set_title(
    "Joy (High Valence, High Arousal) - Left Prefrontal Desync",
    fontsize=10,
    fontweight="bold"
)

axes[1].imshow(imgSad)
axes[1].axis("off")
axes[1].set_title(
    "Sadness (Low Valence, Low Arousal) - Occipital Alpha & Frontal Theta",
    fontsize=10,
    fontweight="bold"
)

plt.tight_layout()
plt.show()
```




Figure 1: 5-band scalp spectral power distributions (Delta, Theta, Alpha, Beta, Gamma) comparing contrasting emotional states.

Frontal Alpha Asymmetry (FAA): How Brain Effort Inverts Alpha Power

A key feature in cognitive science is Frontal Alpha Asymmetry ([Davidson 1992](#); [Coan and Allen 2004](#); [Harmon-Jones, Gable, and Peterson 2010](#)). There is an intuitive biological rule called the **Alpha Paradox**: *alpha waves represent an idle or resting brain state*. When a brain region gets active and starts working hard, its synchronized idling breaks apart (cortical desynchronization), causing alpha power to **drop**.

Therefore, if the left frontal cortex is working harder (associated with approach motivation and positive valence):

$$\text{Alpha}_{\text{Left}} \downarrow \implies \ln(\text{Alpha}_{\text{Right}}) - \ln(\text{Alpha}_{\text{Left}}) > 0$$

We calculate FAA between symmetric forehead electrodes AF4 (right) and AF3 (left):

$$\text{FAA} = \ln(P_{\text{AF4},\alpha}) - \ln(P_{\text{AF3},\alpha})$$

```
fig, ax = plt.subplots(figsize=(9, 5.2), dpi=200)
imgFaa = plt.imread(figuresDir / "01_frontal_alpha_asymmetry.png")
ax.imshow(imgFaa)
ax.axis("off")
plt.tight_layout()
plt.show()
```




Figure 2: Ranked Frontal Alpha Asymmetry (AF4 vs AF3) across the 27 Cowen & Keltner emotion categories (mean $\pm$ SEM).

As shown in Figure 2, positive approach-oriented emotions (Admiration, Craving, Joy, Triumph) have positive FAA values (left-hemisphere dominance), while negative avoidance-oriented emotions (Horror, Fear, Disgust, Sadness) have negative FAA values (right-hemisphere dominance). This confirms our extracted features are tracking real neural dynamics.

Part 2: Manifold Geometry & Representational Similarity Analysis (RSA)

Manifold Unfolding: Diffusion Maps vs. Linear PCA

Our extracted feature space has $D = 462$ dimensions. But we know from geometric deep learning that biological networks don't fill out high-dimensional Euclidean space uniformly; they live on low-dimensional, curved submanifolds (Moon et al. 2019).

We compared standard linear Principal Component Analysis (PCA) against a non-linear **Diffusion Map (Laplacian Eigenmap)** (Belkin and Niyogi 2003): 1. We construct a Gaussian affinity graph across emotion centroids: $A_{ij} = \exp\left(-\frac{\|\mathbf{c}_i - \mathbf{c}_j\|^2}{2\sigma^2}\right)$. 2. We compute degree matrix $D_{ii} = \sum_j A_{ij}$ and construct the row-stochastic Markov random-walk transition matrix $P = D^{-1}A$. 3. We take the leading non-trivial eigenvectors of P to embed the centroids into 2D diffusion coordinates.

```

fig, axes = plt.subplots(1, 2, figsize=(11, 4.5), dpi=200)

imgDiff = plt.imread(figuresDir / "02_manifold_diffusion_quadrants.png")
imgPca = plt.imread(figuresDir / "02_manifold_pca_quadrants.png")

axes[0].imshow(imgDiff)
axes[0].axis("off")
axes[0].set_title("Diffusion Map (Markov Affinity Graph)", fontsize=10, fontweight="bold")

axes[1].imshow(imgPca)
axes[1].axis("off")
axes[1].set_title("PCA (Linear Orthogonal Projection)", fontsize=10, fontweight="bold")

plt.tight_layout()
plt.show()

```




Figure 3: 2D affective embeddings: (Left) Spectral Diffusion Map preserves non-linear clustering, while (Right) Linear PCA collapses into an amorphous blob.

As seen in Figure 3, linear PCA simply squashes the 27 emotions into an unstructured ellipse. In contrast, the Diffusion Map naturally unfolds the manifold: High Valence High Arousal states (Joy, Excitement) branch out smoothly from Low Valence Low Arousal states (Boredom, Sadness).

How Many Dimensions Are We Actually Dealing With? (Two-NN Algorithm)

If our measurement vector has 462 dimensions, how many true degrees of freedom does the affective brain actually need?

Linear PCA suggests we need 5 components for 80% variance and 14 components for 95% variance. But PCA assumes the data lives on flat hyperplanes. To estimate the *intrinsic dimensionality* non-parametrically without assuming linearity, we implemented the **Two-NN algorithm** (Facco et al. 2017).

For every data point i , Two-NN computes the ratio of distance to its 2nd nearest neighbor (r_2) versus its 1st nearest neighbor (r_1):

$$\mu_i = \frac{r_{2,i}}{r_{1,i}} \geq 1$$

Under local Poisson point process assumptions, the cumulative distribution follows a Pareto law: $F(\mu) = 1 - \mu^{-d}$. Fitting this line yields:

$$\hat{d}_{\text{Two-NN}} = 13.42$$

This proves that **the 462-dimensional measurement space collapses onto an intrinsic manifold of only ~13 dimensions**. Roughly 97% of the 462 features are redundant or correlated noise.

Representational Similarity Analysis (RSA): Do Brain Distances Match Psychology?

Next, we wanted to see if the geometry of electrical brain waves matches human psychological perception (Russell 1980; Cowen and Keltner 2017): - **Brain RDM**: The 27×27 pairwise Euclidean distance matrix between emotion centroids in 462D feature space. - **Psychological RDM**: The 27×27 pairwise distance matrix between normative Valence-Arousal coordinates.

We compared both distance matrices using Pearson correlation and Spearman rank correlation across all unique emotion pairs:

```
rsaTable = pd.DataFrame([
    {
        "Metric": "Pearson Correlation (r)",
        "Coeff": f"{rsaR:.4f}",
        "p-value": f"{rsaP:.5f}",
        "Significance": "p < 0.01 (Statistically Significant)",
    },
    {
        "Metric": "Spearman Rank (rho)",
        "Coeff": f"{rsaRho:.4f}",
        "p-value": f"{manifoldMetrics['rsa']['spearman_p']:.5f}",
        "Significance": "p < 0.01 (Statistically Significant)",
    },
])

from IPython.display import Markdown
display(Markdown(rsaTable.to_markdown(index=False)))
```

Table 1: Representational Similarity Analysis (RSA): Correlation between 27x27 Neural EEG Distances and Psychological Circumplex Distances.

Metric	Coeff	p-value	Significance
Pearson Correlation (r)	0.142	0.0077	p < 0.01 (Statistically Significant)
Spearman Rank (rho)	0.1431	0.00725	p < 0.01 (Statistically Significant)

The correlation is positive and statistically significant ($p < 0.01$), proving that when two emotions feel similar to a human, their multi-electrode brain wave patterns are measurably closer in state space.

```
fig, ax = plt.subplots(figsize=(9, 4.5), dpi=200)
imgDendro = plt.imread(figuresDir / "02_emotion_hierarchical_dendrogram.png")
ax.imshow(imgDendro)
ax.axis("off")
plt.tight_layout()
plt.show()
```




Figure 4: Ward’s hierarchical agglomerative clustering dendrogram of the 27 emotion centroids in EEG feature space.

Ward’s clustering tree (Figure 4) groups the 27 emotions into intuitive, biologically grounded subtrees: an avoidance/threat cluster (Fear, Horror, Disgust), an energized approach cluster (Joy, Excitement, Triumph), and a low-energy quiescent cluster (Calmness, Relief, Boredom).

Part 3: Subject-Independent Decoding Benchmark

Stopping Data Leakage: The Leave-Subjects-Out (LSO) Protocol

In many published EEG papers, authors randomly shuffle all trials into 80/20 train/test splits. **This is a huge methodological mistake.** If trial #1 from Subject 5 is in train, and trial #2 from Subject 5 is in test, the classifier doesn’t learn emotion; it memorizes Subject 5’s skull bone thickness, skin conductance, and hair resistance.

To guarantee zero identity leakage, we implemented strict **Leave-Subjects-Out** evaluation via **GroupKFold**:
- **Total Dataset**: 88 unique participants (11,235 total trials) - **Training Set**: 74 subjects (9,438 trials) - **Holdout Test Set**: 14 completely unseen subjects (1,797 trials) - **StandardScaler**: Fit *only* on the 74 training subjects, then applied to the 14 test subjects.

```
quadResults = benchmarkMetrics["task_a_quadrants"]
cowenResults = benchmarkMetrics["task_b_cowen27"]

summaryRecords = []
for q, c in zip(quadResults, cowenResults):
    summaryRecords.append({
        "Model": q["model"],
        "Quad Acc": f"{q['accuracy']*100:.1f}%",
        "Quad F1": f"{q['macro_f1']*100:.1f}%",
        "Top-1": f"{c['accuracy']*100:.1f}%",
        "Top-5": f"{c.get('top_5_accuracy', float('nan'))*100:.1f}%",
        "Train (s)": f"{q['fit_time_s']:.2f}s",
```

```

    "Infer (ms)": f"{q['inference_time_s']*1000:.1f}ms",
  })

dfSummary = pd.DataFrame(summaryRecords)
display(Markdown(dfSummary.to_markdown(index=False)))

```

Table 2: Leave-Subjects-Out (LSO) Performance Benchmark across 6 Model Families on Unseen Participants.

Model	Quad Acc	Quad F1	Top-1	Top-5	Train (s)	Infer (ms)
RandomForest	56.6%	43.4%	30.8%	79.8%	1.37s	27.0ms
MLP_NeuralNet	50.2%	43.2%	27.9%	80.3%	7.05s	4.7ms
ExtraTrees	53.8%	40.9%	26.0%	74.4%	0.17s	27.8ms
Ridge	41.9%	31.5%	12.3%	48.0%	0.06s	2.0ms
LogisticRegression	42.1%	31.4%	10.5%	41.7%	1.67s	1.2ms
NearestCentroid_Co-sine	27.6%	26.0%	5.4%	26.9%	0.02s	7.8ms

What the Benchmarks Show

1. **Circumplex Quadrant Decoding:** Random Forest achieved the highest generalization accuracy (**56.64%** vs. 25.0% random chance), followed closely by our deep Bottleneck MLP (**50.20%**).
2. **27 Fine-Grained Emotions:** Guessing 27 classes across people the model has never seen before is an extremely hard problem (random chance is only 3.70%). Random Forest hit **30.83% Top-1 accuracy** (nearly 5x chance!) and **79.83% Top-5 accuracy**. The Bottleneck MLP reached **80.28% Top-5 accuracy**, meaning in over half the trials, the true emotion is inside the model’s top 5 ranked guesses.
3. **Linear Baselines vs. Deep Models:** Linear Ridge and Logistic Regression capped out at ~12% on the 27-class task. The non-linear models beat them by over 18%, showing that multi-electrode emotion interactions cannot be captured with simple hyperplanes.

```

fig, ax = plt.subplots(figsize=(8, 6.5), dpi=200)
imgCm = plt.imread(figuresDir / "03_confusion_matrix_cowen27.png")
ax.imshow(imgCm)
ax.axis("off")
plt.tight_layout()
plt.show()

```




Figure 5: Confusion matrix for 27 fine-grained emotion decoding under subject-independent evaluation (Random Forest).

Looking at the confusion matrix (Figure 5), the classifier's mistakes make total sense: it confuses semantically close neighbors (like Adoration with Love, or Horror with Fear), rather than confusing opposites (like Joy with Disgust).

Feature Importance: What Electrodes and Frequencies Actually Mattered?

```
fig, ax = plt.subplots(figsize=(9, 4.5), dpi=200)
imgImp = plt.imread(figuresDir / "03_feature_importance_bands_lobes.png")
ax.imshow(imgImp)
ax.axis("off")
plt.tight_layout()
plt.show()
```




Figure 6: Random Forest Gini feature importance grouped by (Left) Frequency Band and (Right) Anatomical Cortical Lobe.

Breaking down Gini feature importance (Figure 6): - **Frequency Bands:** Beta (13–30 Hz) and Alpha (8–13 Hz) together account for over 45% of total tree split importance, followed by the Beta/Alpha ratio. - **Electrode Regions:** Frontal and Prefrontal electrodes account for over 52% of total feature importance, validating that forehead electrodes pick up the strongest cognitive appraisal signals.

Part 4: Generative Modeling with Covariance Shrinkage

The Matrix Singularity Problem in 462 Dimensions

Collecting labeled EEG data from real human participants is expensive, time-consuming, and uncomfortable for subjects. A natural machine learning goal is to build generative models that can synthesize realistic brain wave trials:

$$p(\mathbf{x} \mid y = c) = \mathcal{N}(\mu_c, \Sigma_c)$$

However, in our setup, the feature dimension ($D = 462$) is larger than the number of trials per emotion class. When $D > N$, the standard sample covariance matrix:

$$\mathbf{S}_c = \frac{1}{N-1}(\mathbf{X}_c - \mu_c)^\top(\mathbf{X}_c - \mu_c)$$

is mathematically **singular** (determinant = 0, non-invertible, and full of noisy eigenvalues). You cannot reliably sample from a singular covariance matrix.

Ledoit-Wolf Shrinkage

To solve this, we applied the **Ledoit-Wolf shrinkage estimator** (Ledoit and Wolf 2004). It pulls the noisy empirical covariance toward a well-conditioned identity matrix target:

$$\Sigma_{\text{LW}} = (1 - \rho^*)\mathbf{S}_c + \rho^*\nu\mathbf{I}$$

where $\nu = \frac{1}{p} \text{Tr}(\mathbf{S}_c)$ is the mean variance and $\rho^* \in [0, 1]$ is the analytically optimal shrinkage intensity that minimizes Frobenius loss. This guarantees Σ_{LW} is strictly positive definite and invertible.

```

fig, ax = plt.subplots(figsize=(9, 4.5), dpi=200)
imgGen = plt.imread(figuresDir / "04_generative_feature_fidelity.png")
ax.imshow(imgGen)
ax.axis("off")
plt.tight_layout()
plt.show()

```




Figure 7: Synthetic vs. Real EEG Feature Distributions: (Left) Mean feature vector correlation and (Right) Pairwise covariance spectrum preservation.

Measuring Quality via Fréchet Distance

To verify that the synthetic trials are realistic, we calculated the **Fréchet Distance** (the Wasserstein-2 distance between real and synthetic Gaussian distributions):

$$d_F^2 = \|\mu_r - \mu_s\|_2^2 + \text{Tr}(\Sigma_r + \Sigma_s - 2(\Sigma_r \Sigma_s)^{1/2})$$

- **Sadness:** $d_F = 116.88$
- **Anger:** $d_F = 401.60$
- **Calmness:** $d_F = 1263.00$
- **Joy:** $d_F = 2405.31$

The mean feature vectors achieved a Pearson correlation of $r > 0.98$ with real human data (Figure 7). This shows that regularized Gaussian sampling gives us a fast, stable generative pipeline that avoids the mode collapse and training instability of GANs.

Part 5: Research Directions for Dr. Bhaskar's Lab

As a computer science student interested in geometric ML and computational biology, these findings open up exciting directions that directly connect to **Professor Dhananjay Bhaskar's research group** at UW-Madison:

1. Topological Data Analysis (TDA) on Multi-Channel Time-Series

- **Persistent Homology of EEG Epochs:** Instead of just computing static power bands, we can build sliding-window Vietoris-Rips complexes directly on the raw 14-channel time-series to track 0D and 1D topological persistence barcodes (Betti numbers β_0, β_1) as emotions transition.
- **Topological Loss Functions:** We can add a persistent homology penalty to our Bottleneck MLP loss function to force latent representations to preserve global loop structures.

2. Geometric Deep Learning on Covariance Manifolds

- **Riemannian Manifold Neural Networks:** EEG spatial covariance matrices $\mathbf{P} = \frac{1}{T}\mathbf{X}\mathbf{X}^\top$ live on the manifold of Symmetric Positive-Definite (SPD) matrices $\mathcal{S}_{++}^p$. Instead of flattening matrices into Euclidean vectors, we can use Riemannian geometry (like the Log-Euclidean Metric or Affine-Invariant Riemannian Metric) with SPDNet architectures.
- **Graph Neural Networks (GNNs):** Treating the 14 electrodes as nodes in a dynamic graph where edge weights represent time-varying coherence or phase-locking values.

3. Non-Linear Trajectory Inference

- **PHATE & Diffusion Trajectories:** We can use tools like PHATE (Moon et al. 2019) to trace continuous trajectories between emotions, identifying the exact branch points where mixed emotional states (like bittersweet or nostalgic feelings) bifurcate into distinct categorical states.

References

- Belkin, Mikhail, and Partha Niyogi. 2003. “Laplacian Eigenmaps for Dimensionality Reduction and Data Representation.” *Neural Computation* 15 (6): 1373–96. <https://doi.org/10.1162/089976603321780317>.
- Coan, James A, and John JB Allen. 2004. “Frontal EEG Asymmetry as a Moderator and Mediator of Emotion.” *Biological Psychology* 67 (1-2): 7–49. <https://doi.org/10.1016/j.biopsycho.2004.03.002>.
- Cowen, Alan S, and Dacher Keltner. 2017. “Self-Report Captures 27 Distinct Categories of Emotion Bridged by Continuous Gradients.” *Proceedings of the National Academy of Sciences* 114 (38): E7900–7909. <https://doi.org/10.1073/pnas.1702247114>.
- Davidson, Richard J. 1992. “Anterior Cerebral Asymmetry and the Nature of Emotion.” *Brain and Cognition* 20 (1): 125–51. [https://doi.org/10.1016/0278-2626\(92\)90014-M](https://doi.org/10.1016/0278-2626(92)90014-M).
- Facco, Elena, Maria d’Errico, Alex Rodriguez, and Alessandro Laio. 2017. “Estimating the Intrinsic Dimension of Datasets by a Minimal Neighborhood Information.” *Scientific Reports* 7 (1): 12140. <https://doi.org/10.1038/s41598-017-11873-y>.
- Harmon-Jones, Eddie, Philip A Gable, and Cindy K Peterson. 2010. “The Role of Asymmetric Frontal Cortical Activity in Emotion-Related Phenomena: A Review and Update.” *Biological Psychology* 84 (3): 451–62. <https://doi.org/10.1016/j.biopsycho.2009.08.010>.
- Kriegeskorte, Nikolaus, Marieke Mur, and Peter A Bandettini. 2008. “Representational Similarity Analysis—Connecting the Branches of Systems Neuroscience.” *Frontiers in Systems Neuroscience* 2: 4. <https://doi.org/10.3389/neuro.06.004.2008>.
- Ledoit, Olivier, and Michael Wolf. 2004. “A Well-Conditioned Estimator for Large-Dimensional Covariance Matrices.” *Journal of Multivariate Analysis* 88 (2): 365–411. [https://doi.org/10.1016/S0047-259X\(03\)00096-4](https://doi.org/10.1016/S0047-259X(03)00096-4).
- Moon, Kevin R, David van Dijk, Zheng Wang, Scott Gigante, Daniel B Burkhardt, William S Chen, Kristina Yim, et al. 2019. “Visualizing Structure and Transitions in High-Dimensional Biological Data.” *Nature Biotechnology* 37 (12): 1482–92. <https://doi.org/10.1038/s41587-019-0336-3>.
- Phuong, Hoang-Tung, Eun-Taek Im, Min-Seok Oh, and Gwang-Yong Gim. 2025. “EEGEmotions-27: A Large-Scale EEG Dataset Annotated with 27 Fine-Grained Emotion Labels.” *IEEE Access* 13: 176915–32. <https://doi.org/10.1109/ACCESS.2025.3620677>.

Russell, James A. 1980. "A Circumplex Model of Affect." *Journal of Personality and Social Psychology* 39 (6): 1161–78. <https://doi.org/10.1037/h0077714>.